



HORTONWORKS®



What is Hortonworks?



The Sandbox by **Hortonworks** is a straightforward, pre-configured, learning environment that contains the latest developments from Apache Hadoop, specifically the Hortonworks Data Platform (HDP).

The Sandbox comes packaged in a virtual environment that can run in the cloud or on your machine.

Sandbox also offers a data-in-motion framework for IoT solutions called Hortonworks Data Flow (HDF).

Hortonworks offers a way to use Hadoop Tools connecting to a Virtual Machine in SSH for command lines interfaces. Numerous web interfaces are also available.

Download the **HDP** Sandbox. This Sandbox makes it easy to get started with Apache Hadoop, Apache Spark, Apache Hive, Apache HBase, Druid and Data Analytics Studio (DAS). Choose the VirtualBox installation type.

Hortonworks Data Platform (HDP[®]) on Hortonworks Sandbox

The HDP Sandbox makes it easy to get started with Apache Hadoop, Apache Spark, Apache Hive, Apache HBase, Druid and Data Analytics Studio (DAS).

Get Started Now

CHOOSE INSTALLATION TYPE



Virtualbox

VMWARE

Docker

Fill in the form and download the **version 2.6.5** of the Sandbox. The sandbox requires around 15 Go of space.

Once this is done (the download might take a long time) :

- Open VirtualBox
- Click on “New”
- Go in “File -> Import a Virtual Machine”.
- Select the image of the Sandbox you just downloaded.
- Start the VM

Oracle VM VirtualBox - Gestionnaire de machines

Outils

Nouvelle

Configuration

Oublier

Démarrer

Linux

Éteinte

Hortonworks Sandbox HDP 2.6.5

Éteinte

cloudera-quickstart-vm-5.13.0-0-...

Éteinte

PostgreSQL16

Sauvegardée

cassandra_spark_cluster_5_1

Éteinte

Linux_VivaData

Éteinte

Hortonworks Sandbox HDP ...

Éteinte

Général

Nom : Hortonworks Sandbox HDP 2.6.5 1

Système d'exploitation : Red Hat (64-bit)

Emplacement du fichier de paramètres : /Users/maelfabien/VirtualBox VMs/Hortonworks Sandbox HDP 2.6.5 1

Système

Mémoire vive : 8192 Mo

Processeurs : 4

Ordre d'amorçage : Disque dur, Optique

Accélération : VT-x/AMD-V , Pagination imbriquée, PAE/NX , Paravirtualisation KVM

Affichage

Mémoire vidéo : 8 Mo

Contrôleur graphique : VBoxVGA

Port serveur bureau distant : 5905

Enregistrement : Désactivé

Stockage

Contrôleur : IDE Controller

Maître primaire IDE : Hortonworks Sandbox HDP 2.6.5-disk001.vdi (Normal, 117,19 Gio)

Son

Pilote hôte : CoreAudio

Contrôleur : ICH AC97

Réseau

Interface 1: Intel PRO/1000 MT Desktop (NAT)

USB

Prévisualisation

Hortonworks Sandbox
HDP 2.6.5 1

The first boot takes a while, so time for a break!

```
====  
First boot detected :)  
Please wait while the Hortonworks Sandbox is extracted and deployed.  
This process will only occur once.  
Grab a drink - this may take a few minutes.  
====  
  
While you wait, read up on interesting data articles here:  
https://community.hortonworks.com/kb/list.html  
  
Extracting and loading the Hortonworks Sandbox...
```


Access the Sandbox

The application is now ready to be used :

```
Hortonworks HDP Sandbox  
https://hortonworks.com/products/sandbox
```

```
To quickly get started with the Hortonworks Sandbox, follow this tutorial:  
https://hortonworks.com/tutorial/hadoop-tutorial-getting-started-with-hdp/
```

```
To initiate your Hortonworks Sandbox session, open a browser to this address:
```

```
For VirtualBox:
```

```
Welcome screen: http://localhost:1080
```

```
SSH: http://localhost:4200
```

```
For VMWare:
```

```
Welcome screen: http://10.0.2.15:1080
```

```
SSH: http://10.0.2.15:4200
```

User

Once started, the Sandbox is accessible from your local computer, in SSH, on the Port 2222. We will be using the username: `raj_ops` and password `raj_ops` as it is pre-registered.

There are many pre-configured users for the HortonWorks Sandbox, including :

User	Password	Role	Service	Ambari Authorization
Admin	refer to Admin Password Reset	Ambari Admin	Ambari	Ambari Admin
raj_ops	raj_ops	Hadoop Warehouse Operator	Hive/Tez, Ranger, Falcon, Knox, Sqoop, Oozie, Flume, Zookeeper	Cluster Administrator
maria_dev	maria_dev	Spark and SQL Developer	Hive, Zeppelin, MapReduce/Tez/Spark, Pig, Solr, HBase/Phoenix, Sqoop, NiFi, Storm, Kafka, Flume	Service Operator
holger_gov	holger_gov	Intendant des données	Atlas	Service Administrator
amy_ds	amy_ds	Data Scientist	Spark, Hive, R, Python, Scala	Service Operator

SSH

Launch your terminal and access the Sandbox using SSH :

```
ssh maria_dev@localhost -p 2222
```

If you are asked whether you'd like to permanently add 2222 to the list of known hosts, say Yes.

Splash Page

We should be able to access a graphical view of the services available on a Hadoop cluster. It's called the VirtualBox Splash Page, and it can be accessed on :

<http://localhost:8080>

You will be asked to log-in :



Sign in

Username

Password

Sign in



- ✓ HDFS
- ✓ YARN
- ✓ MapReduce2
- Tez
- ✓ Hive
- HBase
- Pig
- Sqoop
- ✓ Oozie
- ✓ ZooKeeper
- Falcon
- Storm
- ✓ Flume
- Ambari Infra
- Atlas
- ✓ Kafka
- Knox
- ✓ Ranger
- ✓ Spark2
- ✓ Zeppelin
- Notebook
- Druid
- Slider
- Superset

Actions

Metrics

Heatmaps

Config History

Metric Actions

Last 1 hour

HDFS Disk Usage



DataNodes Live

1/1

HDFS Links

NameNode
Secondary NameNode
1 DataNodes

More...

Memory Usage

No Data Available

Network Usage

No Data Available

CPU Usage

No Data Available

Cluster Load

No Data Available

NameNode Heap



NameNode RPC

0.90 ms

NameNode CPU WIO

n/a

NameNode Uptime

5.6 hr

HBase Master Heap

n/a

HBase Links

No Active Master
1 RegionServers
n/a

More...

HBase Ave Load

n/a

HBase Master Uptime

n/a

ResourceManager Heap



ResourceManager Uptime

5.6 hr

YARN Memory



NodeManagers Live

1/1

YARN Links

ResourceManager
1 NodeManagers

More...

An aerial photograph of New York City at dusk. The Empire State Building is prominent in the center, with its top illuminated in red and green. The city is densely packed with skyscrapers, many of which have their lights on. The sky is a mix of dark blue and orange from the setting sun. The word "Questions?" is written in large, white, sans-serif font across the middle of the image.

Questions?